REVOLUTION ANALYTICS

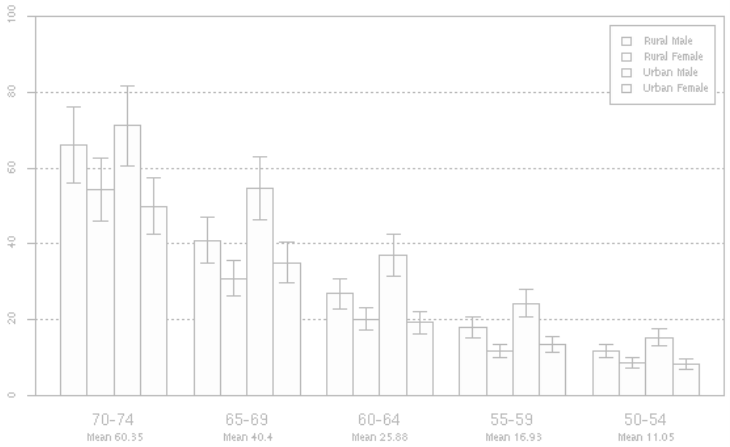

Revolution Analytics

R and Data Science

Joseph B Rickert

September 25, 2014

What is R?

- Most widely used data analysis software
  - Used by 2M+ data scientists, statisticians and analysts
- Most powerful statistical programming language
  - Flexible, extensible and comprehensive for productivity
- Platform for beautiful and unique data visualizations
  - As seen in New York Times, Twitter and Flowing Data
- Thriving open-source community
  - Leading edge of analytics research

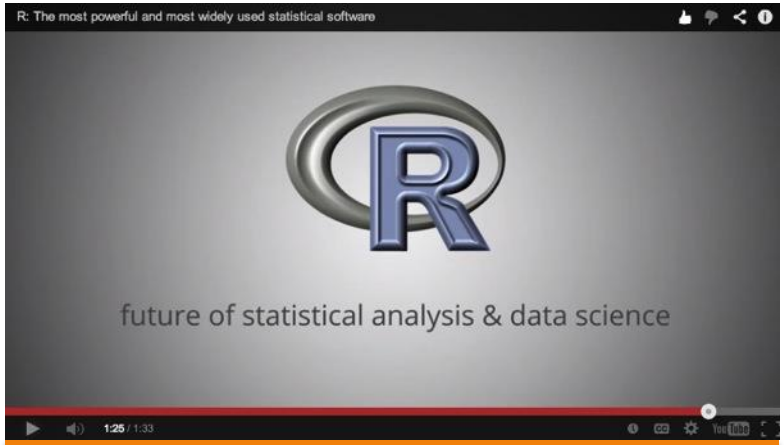

www.revolutionanalytics.com/what-r

REVOLUTION ANALYTICS

OPEN SOURCE R

R's popularity is growing rapidly

R Usage Growth

Rexer Data Miner Survey, 2007-2013

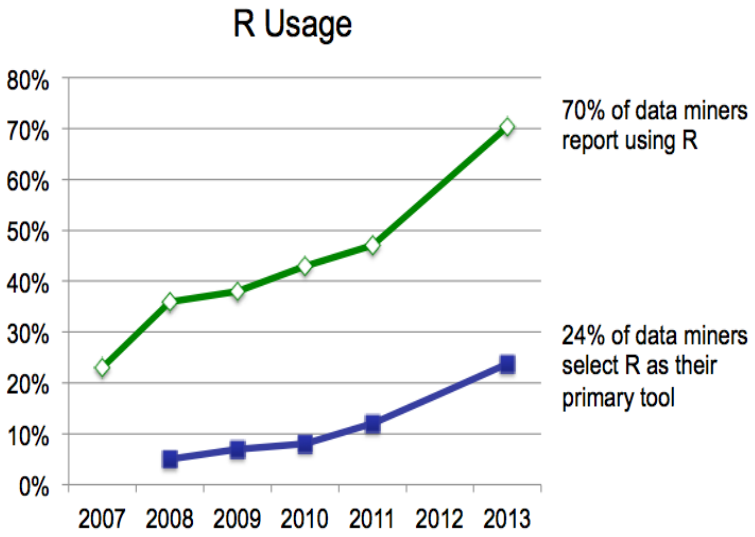

- Rexer Data Miner Survey

Language Popularity

IEEE Spectrum Top Programming Languages

Language Rank	Types	Spectrum Ranking
1. Java		100.0
2. C		99.2
3. C++		95.5
4. Python		93.4
5. C#		92.2
6. PHP		84.6
7. Javascript		84.3
8. Ruby		78.6
9. R		74.0
10. MATLAB		72.6

#9: R

- IEEE Spectrum, July 2014

REVOLUTION ANALYTICS

Poll Question #1

- What are the statistical programming languages/platforms you are most familiar with? (choose all that apply)
  - A) R
  - B) SAS
  - C) SPSS
  - D) KXEN
  - E) Statistica

Tools for Data Science

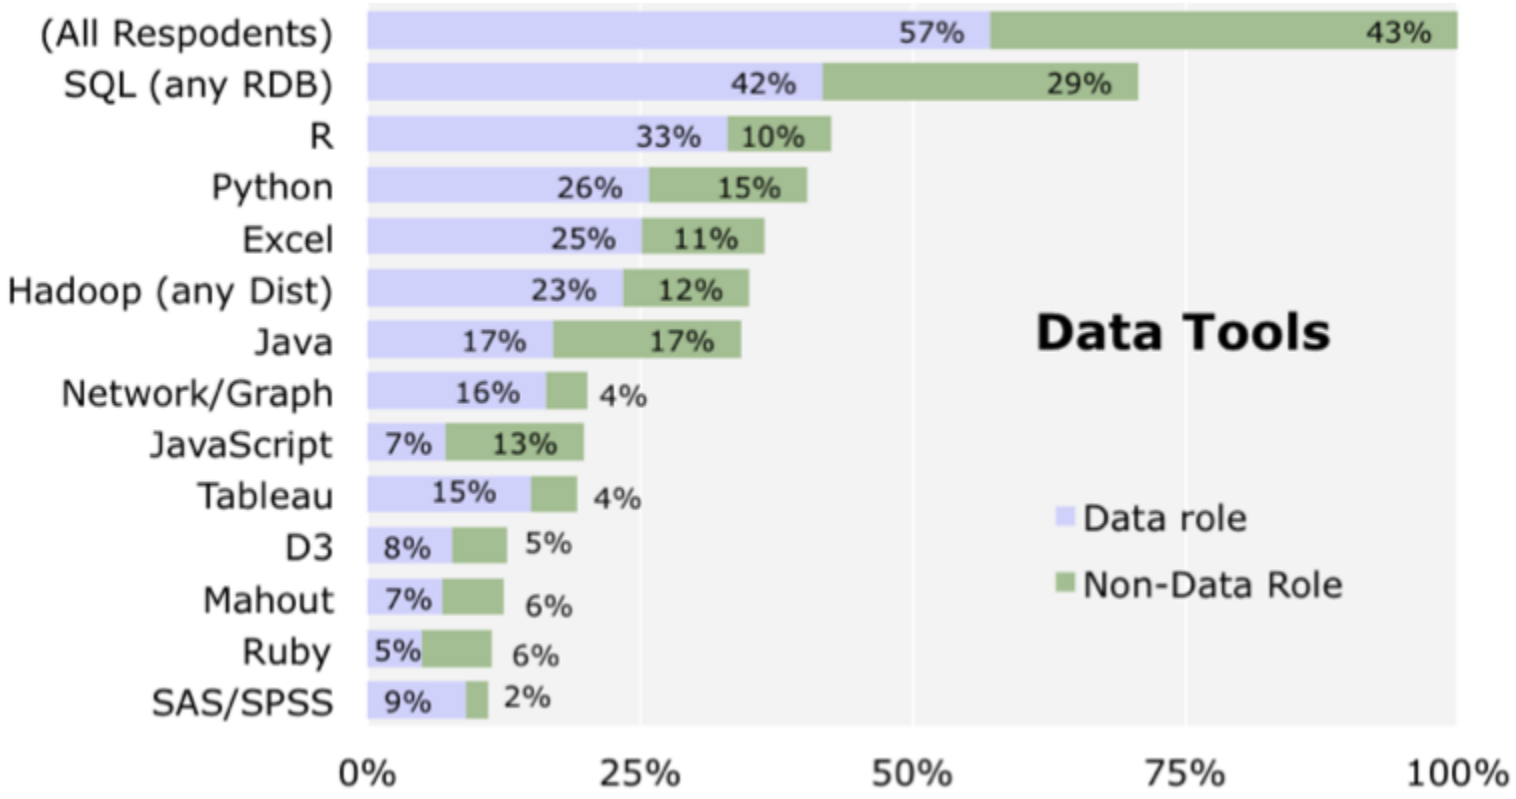

Source: O'Reilly Data Science Survey

R is among the highest-paid IT skills in the US

AVERAGE SALARY FOR High Paying Skills and Experience

SKILL	2013	YR/YR CHANGE
R	$ 115,531	n/a
NoSQL	$ 114,796	1.6%
MapReduce	$ 114,396	n/a
PMBok	$ 112,382	1.3%
Cassandra	$ 112,382	n/a
Omnigraffle	$ 111,039	0.3%
Pig	$ 109,561	n/a
SOA (Service Oriented Architecture)	$ 108,997	-0.5%
Hadoop	$ 108,669	-5.6%
Mongo DB	$ 107,825	-0.4%

Dice Tech Salary Survey, January 2014

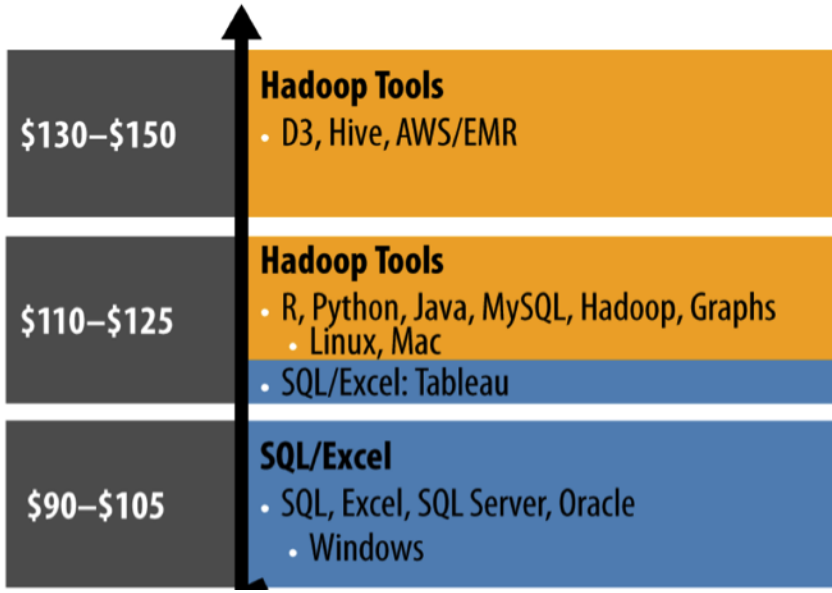

O'Reilly Strata 2013 Data Science Salary Survey

Photo by Ksayer1 on flickr.

Why R for Data Science?

Algorithms

Task Views

```
X <- if (!is.empty.model(mt))
    model.matrix(mt, mf, contrasts)
  else matrix(, NROW(Y), 0L)
  weights <- as.vector(model.weights(mf))
  if (!is.null(weights) && !is.numeric(weights))
    stop("'weights' must be a numeric vector")
  if (!is.null(weights) && any(weights < 0))
    stop("negative weights not allowed")
  offset <- as.vector(model.offset(mf))
  if (!is.null(offset)) {
    if (length(offset) != NROW(Y))
      stop(gettextf("number of offsets is %d should equal %d (number of observations)",
        length(offset), NROW(Y)), domain = NA)
  }
  mustart <- model.extract(mf, "mustart")
  etastart <- model.extract(mf, "etastart")
  fit <- eval(call(if (is.function(method)) "method" else method,
    x = X, y = Y, weights = weights, start = start, etastart = etastart,
    mustart = mustart, offset = offset, family = family,
    control = control, intercept = attr(mt, "intercept") >
      0L))
  if (length(offset) && attr(mt, "intercept") > 0L) {
    fit2 <- eval(call(if (is.function(method)) "method" else method,
      x = X[, "(Intercept)", drop = FALSE], y = Y, weights = weights,
      offset = offset, family = family, control = control,
      intercept = TRUE))
    if (!fit2$converged)
      warning("fitting to calculate the null deviance did not converge -- increase 'maxit'?")
    fit$null.deviance <- fit2$deviance
  }
  if (model)
    fit$model <- mf
  fit$na.action <- attr(mf, "na.action")
  if (x)
    fit$x <- X
  if (!y)
    fit$y <- NULL
  fit <- c(fit, list(call = call, formula = formula, terms = mt,
    data = data, offset = offset, control = control, method = method,
    contrasts = attr(X, "contrasts"), xlevels = .getXlevels(mt,
      mf)))
  class(fit) <- c(fit$class, c("glm", "lm"))
  fit
```

REVOLUTION ANALYTICS

R Growth

Put this astonishing growth in perspective:

- SAS.V 9.3S contains ~ 1,200 commands that are roughly equivalent to R functions
- R packages contain a median of 5 functions
- Therefore R has ~ 36,820 functions
- *During 2013 alone, R added more functions than SAS Institute has written in its entire history!*

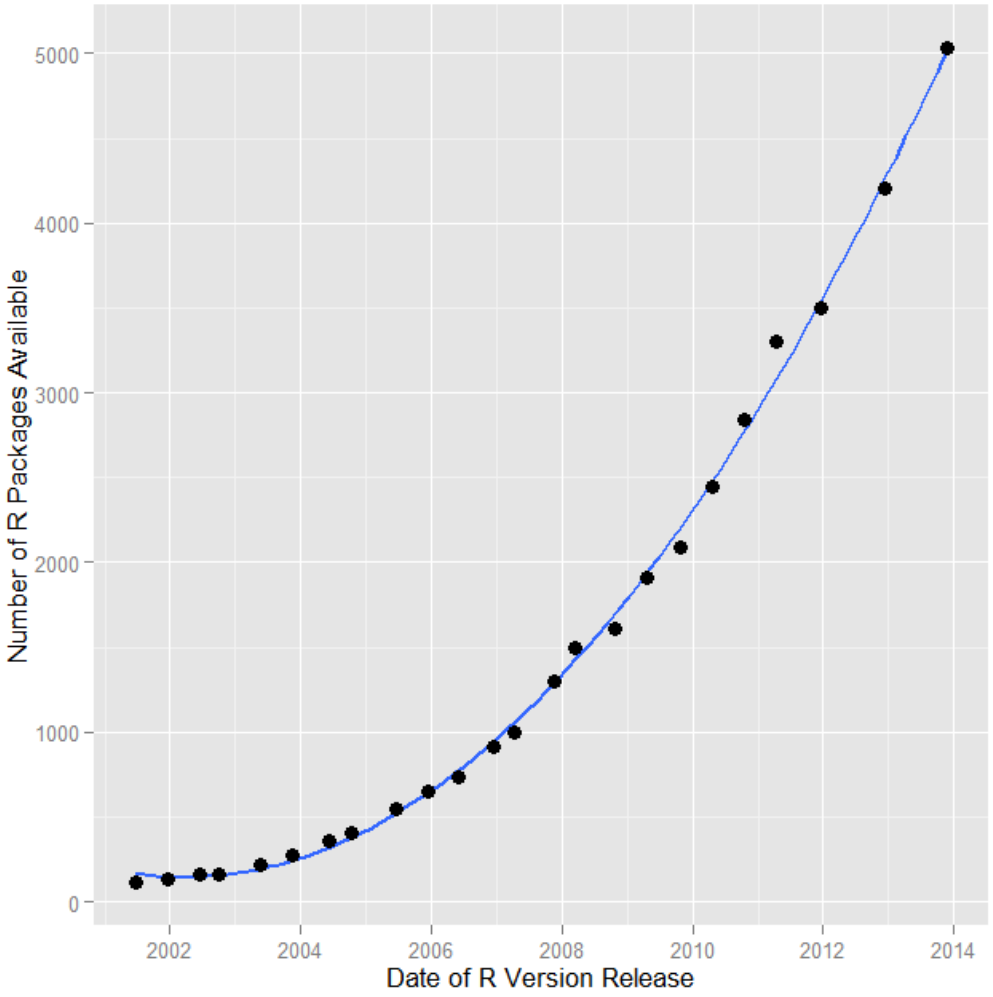

Bob Muenchen

Why R for Data Science?

Visualizations

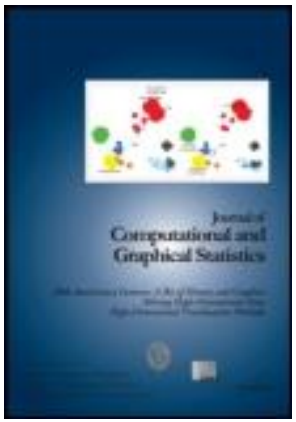

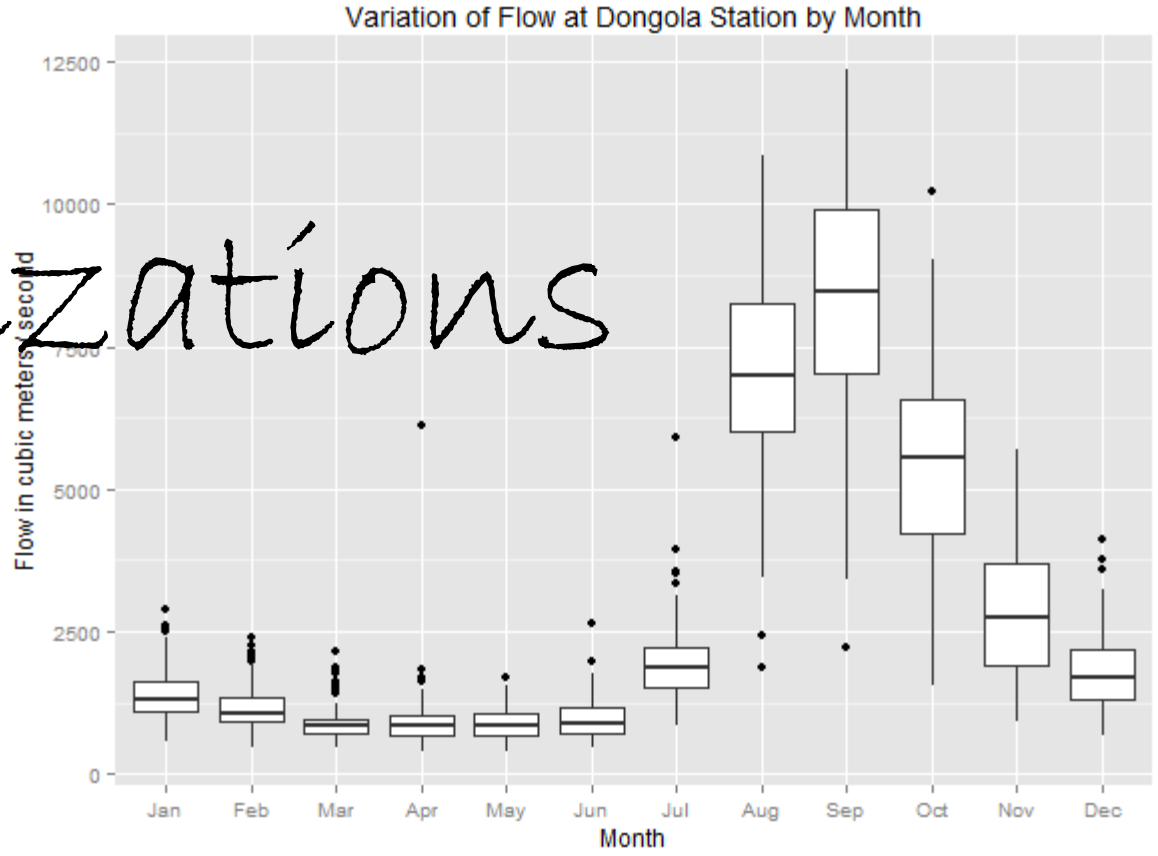

REVOLUTION ANALYTICS

Why R for Data Science?

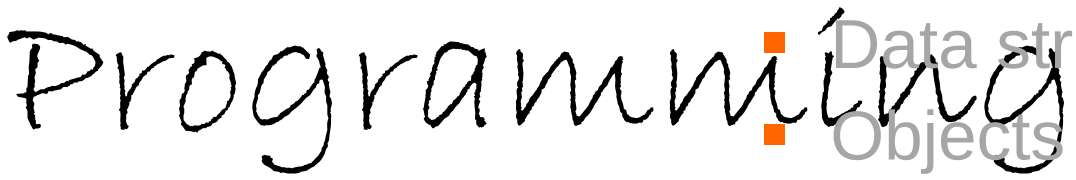

- Scripting
- Functional programming
- Parallel programming
- Data structures
- Objects
- Data Types
- Regular expressions
- Data connections
- Interfaces to other languages

Why R for Data Science?

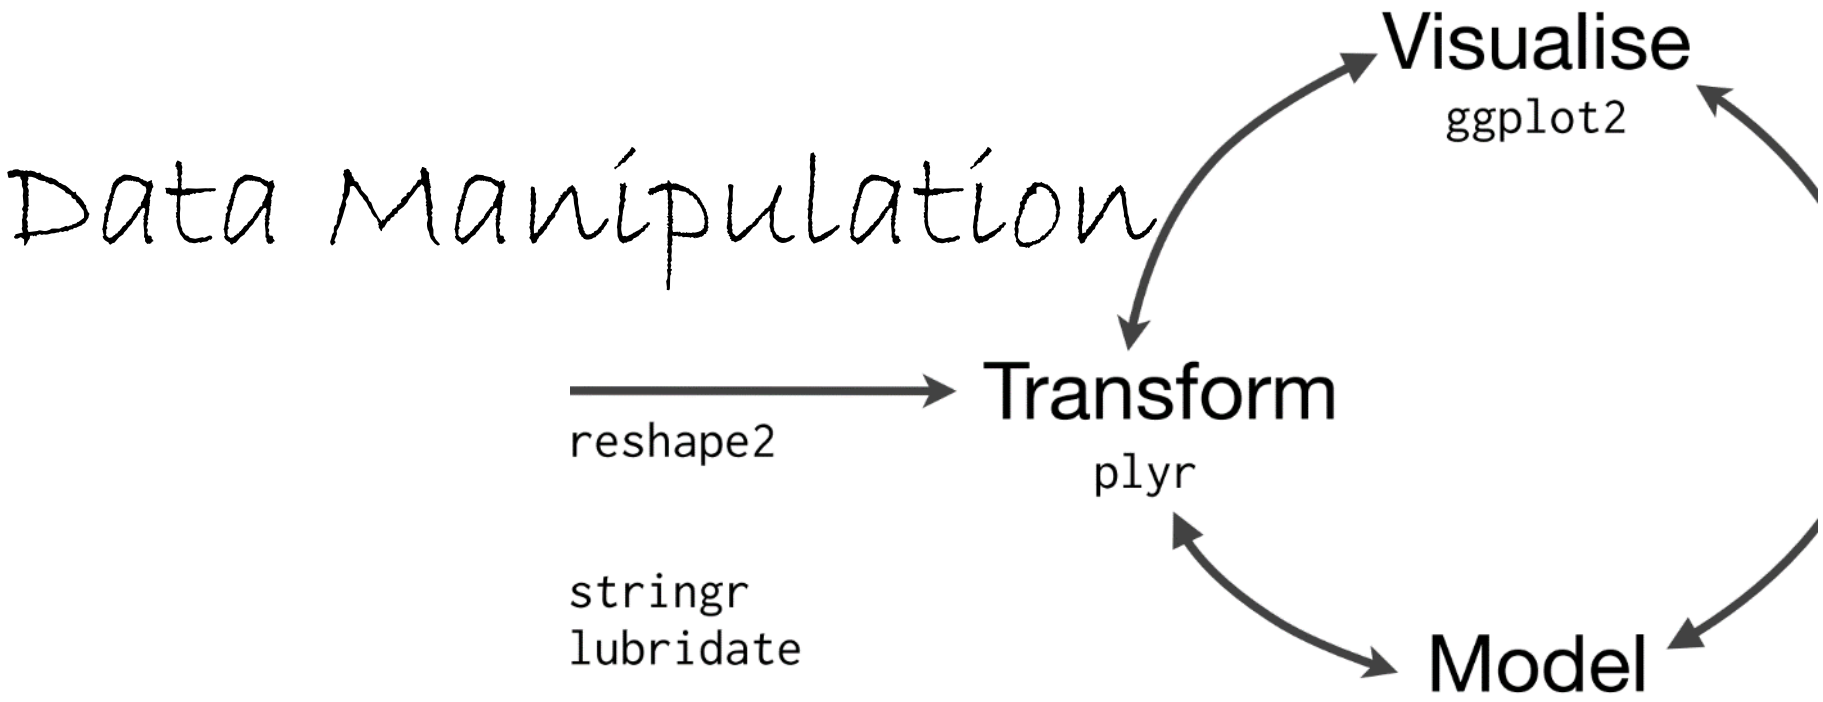

"It's often said that 80% of the effort of analysis is spent just getting the data ready to analyse, the process of data cleaning. Data cleaning is not only a vital first step, but it is often repeated multiple times over the course of an analysis as new problems come to light." Hadley Wickham Tidy Data

Why R for Data Science?

R Integrates

- Web applications
- Internet graphics
  - D3
  - Potly
- Other Languages
  - C, C++
  - Java
- BI Tools
- Data bases
  - SQL
  - MongoDB

Poll Question #2

- What are the data platforms that you are connecting to regularly? (choose all that apply)
  - A) Hadoop
  - B) Spark
  - C) Cloud-based (Azure/AWS/Google)
  - D) Data Warehouses
  - E) Servers (Grid or Cluster)

Why R for Data Science

R Scales

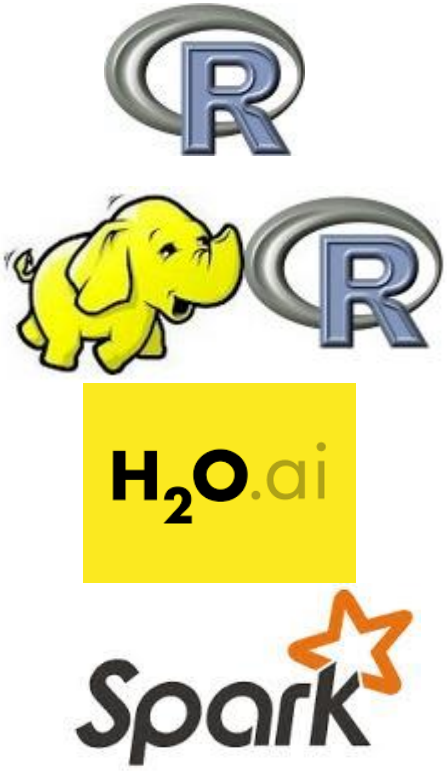

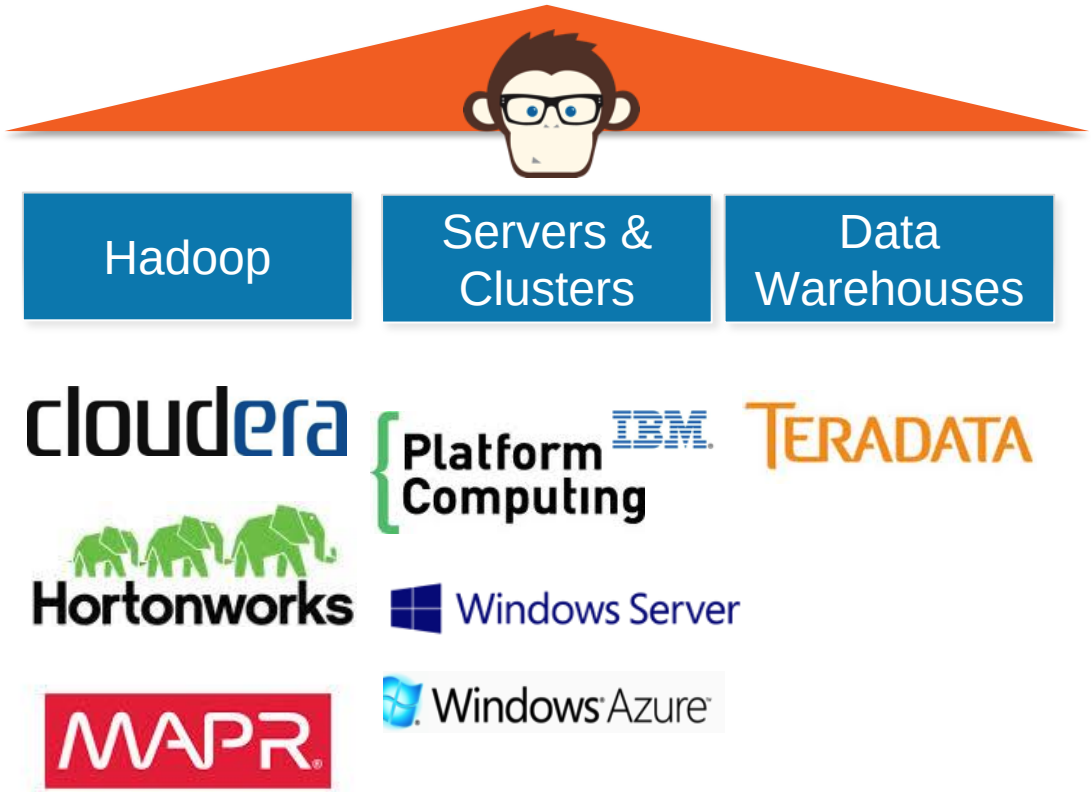

Poll Question #3

- What are the types of models that you are working with most? (choose all that apply)
  - A) Linear models / Regression / GLM
  - B) Decision Trees / Random Forests
  - C) Survival Models
  - D) GBM
  - E) Time Series models

REVOLUTION ANALYTICS

Let's look at some code.

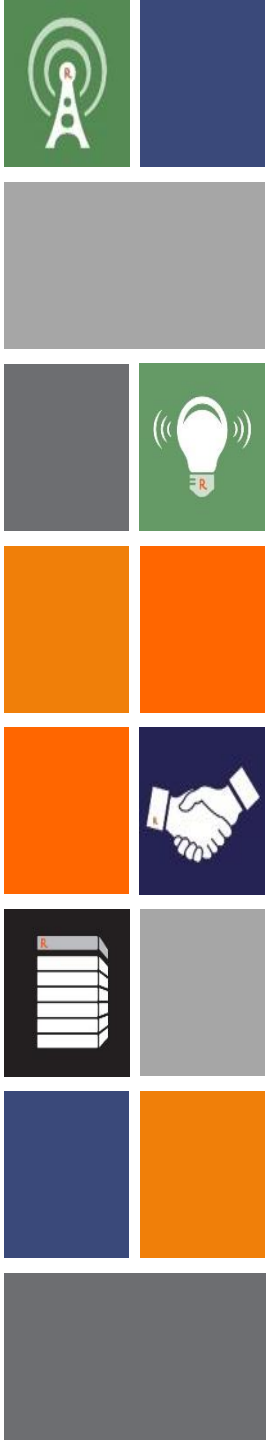

www.revolutionanalytics.com
1.855.GET.REVO
Twitter: @RevolutionR

Why is R Right for Data Science?

- R is open source
- R is a powerful language
  - Data Manipulation
  - Computational Statistics
  - Machine Learning
- R is an innovation engine
- R has a rich and expanding ecosystem

Q&A / Resources

R Code and Markdown Files
https://github.com/joseph-rickert/DataScienceRWebinar

What is R?
revolutionanalytics.com/what-is-r

Companies using R
revolutionanalytics.com/companies-using-r

AcademyR training
revolutionanalytics.com/AcademyR

AcademyR Certification
revolutionanalytics.com/AcademyR-certification

Contact Revolution Analytics
revolutionanalytics.com/contact-us

REVOLUTION ANALYTICS

REVOLUTION ANALYTICS

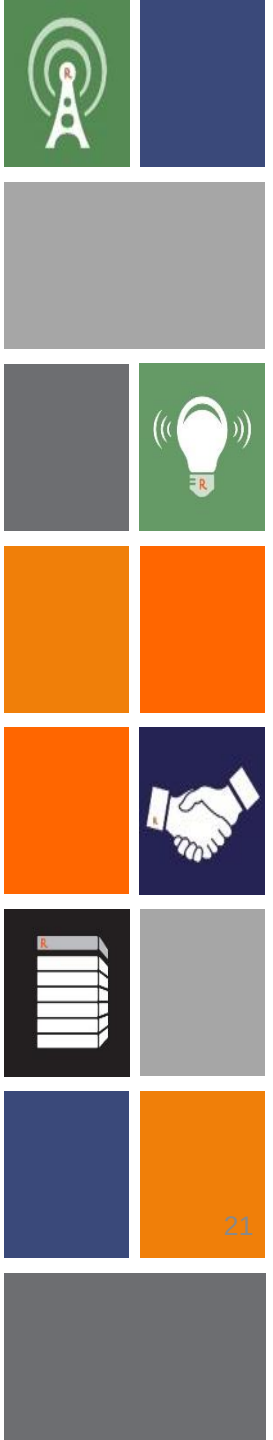

Thank you

Revolution Analytics is the leading commercial provider of software and support for the popular open source R statistics language.

www.revolutionanalytics.com, 1.855.GET.REVO, Twitter: @RevolutionR